Spatial dimension:

We can give every point in our three-dimensional space an *(x, y, z)* coordinate. Since these points are represented as three real numbers, we say that it belongs to the set $\mathbb{R}$. These coordinates are usually not very meaningful, but they encode useful notions of distance and magnitude.

The important thing above is that when we had *three* real-valued numbers, we could represent it in our three-dimensional space. Let's now think about an image. An image is simply a big collection of pixels, with each pixel representing an intensity in some range. Here spatial Dimension is Row and Column and depth dimension is 3 (R, G, B).

LSTM:

LSTM recurrent unit tries to "remember" all the past knowledge that the network is seen so far and to "forget" irrelevant data.

4 gates control how much information have to memorize.

Row LSTM: Hidden state (i, j) = Hidden state (i-1, j-1), Hidden state (i-1, j+1), Hidden state(i-1, j), p(i, j)

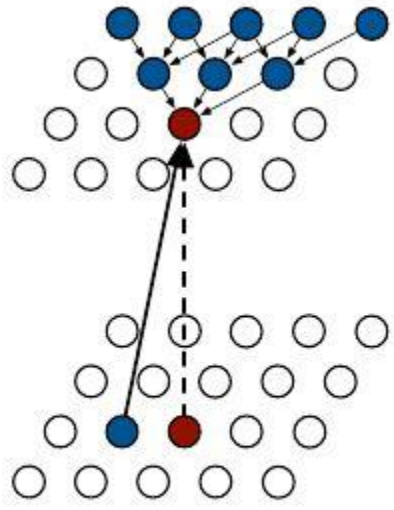

Row LSTM

Bi-Diagonal LSTM: In diagonal BLSTM, the hidden state of a pixel (i, j) depends on pixel (i, j-1) and on pixel (i-1, j). As bidirectional LSTM covers forward and backward dependencies, all the previously generated pixels are included in the 'context' / 'history' used in predicting value for a pixel.

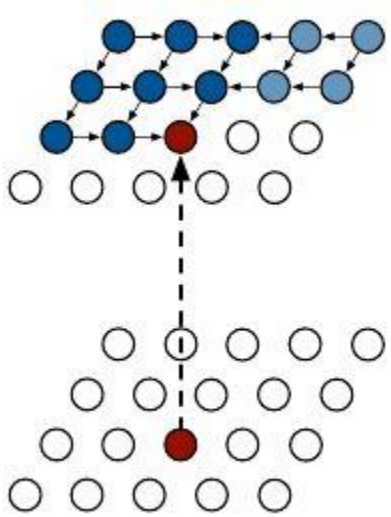

The CIFAR-10 dataset

The CIFAR-10 dataset consists of 60000 32x32 color images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images

Some image generation process:

1. VAE (Variational Auto Encoder): Generate image efficiently and quickly but blurry.

2. GAN (Generative Adversarial Networks): Generate more sharper image than VAE but difficult to optimize.

3. Auto Regressive Model (Ex. Pixel CNN, Pixel RNN): Simple and stable training process. Problem is it is inefficient during sampling.